

```
In [1]: #imported libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: df = pd.read_csv(r'perrin-freres-monthly-champagne-.csv') #Imported Data
df
```

```
Out[2]:
```

	Month	Perrin Freres monthly champagne sales millions ?64-?72
0	1964-01	2815.0
1	1964-02	2672.0
2	1964-03	2755.0
3	1964-04	2721.0
4	1964-05	2946.0
...	...	...
102	1972-07	4298.0
103	1972-08	1413.0
104	1972-09	5877.0
105	NaN	NaN
106	Perrin Freres monthly champagne sales millions...	NaN

107 rows × 2 columns

```
In [3]: df.info() #Checking the Data type and null values
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 107 entries, 0 to 106
Data columns (total 2 columns):
#   Column                                                                 Non-Null Count  Dtype
---  -
0   Month                                                                106 non-null   object
1   Perrin Freres monthly champagne sales millions ?64-?72          105 non-null   float64
dtypes: float64(1), object(1)
memory usage: 1.8+ KB
```

```
In [4]: round(df.describe(),2) #Checking statistical properties
```

Out[4]: **Perrin Freres monthly champagne sales millions ?64-?72**

<b>count</b>	105.00
<b>mean</b>	4761.15
<b>std</b>	2553.50
<b>min</b>	1413.00
<b>25%</b>	3113.00
<b>50%</b>	4217.00
<b>75%</b>	5221.00
<b>max</b>	13916.00

```
In [5]: df.dropna(inplace=True) #Dropping null values
df.isnull().sum() #Checking null values
```

```
Out[5]: Month 0
Perrin Freres monthly champagne sales millions ?64-?72 0
dtype: int64
```

```
In [6]: df['Month'] = pd.to_datetime(df['Month'],format='mixed') #Converting month Dtype in
df.info() #Checking info of the data
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 105 entries, 0 to 104
Data columns (total 2 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Month                                105 non-null    datetim
e64[ns]
1   Perrin Freres monthly champagne sales millions ?64-?72  105 non-null    float64
dtypes: datetime64[ns](1), float64(1)
memory usage: 2.5 KB
```

```
In [7]: df.tail() #shows top 5 bottom rows
```

Out[7]: **Month Perrin Freres monthly champagne sales millions ?64-?72**

<b>100</b>	1972-05-01	4618.0
<b>101</b>	1972-06-01	5312.0
<b>102</b>	1972-07-01	4298.0
<b>103</b>	1972-08-01	1413.0
<b>104</b>	1972-09-01	5877.0

```
In [8]: df.shape #shows shape of the data
```

Out[8]: (105, 2)

```
In [9]: df.columns = ['Month','Sale'] #Changes the aname of the columns
df
```

```
Out[9]:
```

	Month	Sale
0	1964-01-01	2815.0
1	1964-02-01	2672.0
2	1964-03-01	2755.0
3	1964-04-01	2721.0
4	1964-05-01	2946.0
...	...	...
100	1972-05-01	4618.0
101	1972-06-01	5312.0
102	1972-07-01	4298.0
103	1972-08-01	1413.0
104	1972-09-01	5877.0

105 rows × 2 columns

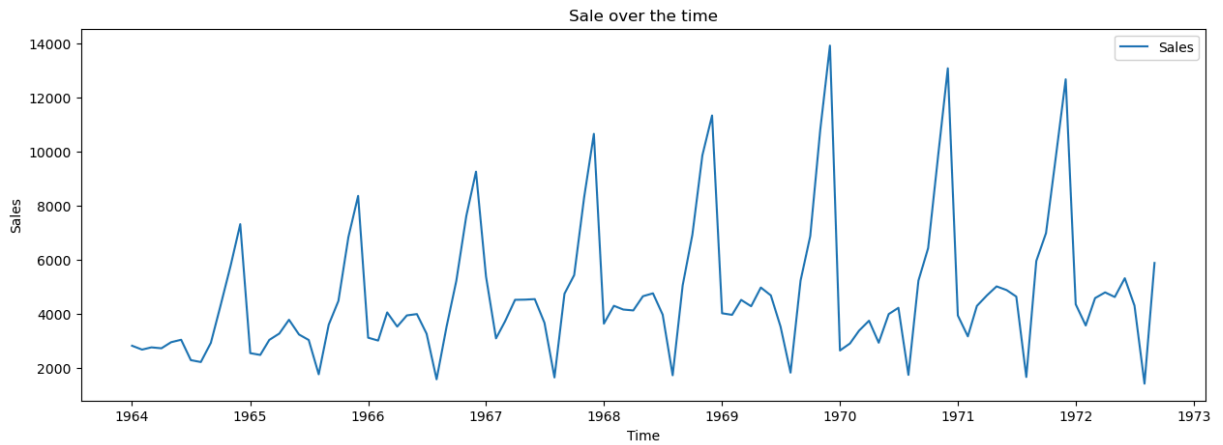
```
In [10]: df.set_index('Month', inplace=True) #Set month columns into index
```

```
In [11]: df.head() #Shows top 5 columns
```

```
Out[11]:
```


	Sale
Month	
1964-01-01	2815.0
1964-02-01	2672.0
1964-03-01	2755.0
1964-04-01	2721.0
1964-05-01	2946.0

```
In [12]: plt.figure(figsize=(15,5))
plt.title('Sale over the time')
plt.plot(df, label = 'Sales')
plt.xlabel('Time')
plt.ylabel('Sales')
plt.legend()
plt.show()
```



```
In [13]: rollmean = df.rolling(window=12).mean()
rollstd = df.rolling(window=12).std()
```

```
plt.figure(figsize=(15,5))
plt.title('Sale over the time')
plt.plot(df, label = 'Original')
plt.plot(rollmean, label = 'Mean')
plt.plot(rollstd, label = 'Std')
plt.legend()
plt.show()
```



```
In [14]: #checking stationarity
from statsmodels.tsa.stattools import adfuller
```

```
In [15]: # H0: Data is not stationary
# H1: Data is stationary

def test_stationarity(x):
    adf_test = adfuller(x)
    print('ADF Test Results')
    print(f'Statistic Value: {adf_test[0]}')
    p_value = adf_test[1]
    print(f'P-Value: {p_value}')
    print('Critical Values:')
    for key, value in adf_test[4].items():
        print(f'    {key}: {value}')
```

```

if p_value < 0.05:
    print('✅ Series is stationary. Reject the null hypothesis (H0).')
else:
    print('❌ Series is NOT stationary. accept the null hypothesis (H0).')

```

In [16]: `test_stationarity(df)`

```

ADF Test Results
Statistic Value: -1.833593056327623
P-Value: 0.363915771660245
Critical Values:
    1%: -3.502704609582561
    5%: -2.8931578098779522
   10%: -2.583636712914788
❌ Series is NOT stationary. accept the null hypothesis (H0).

```

In [17]: `from statsmodels.tsa.seasonal import seasonal_decompose`  
`seasonl_additive = seasonal_decompose(df, period=12, model = 'addative')`

In [18]: `trend = seasonl_additive.trend`  
`seasonal = seasonl_additive.seasonal`  
`residual = seasonl_additive.resid`

In [19]: `plt.figure(figsize=(15,10))`  
`plt.title('Sale over the time')`  
  
`plt.subplot(411)`  
`plt.plot(df, label = 'Orignal')`  
`plt.legend()`  
  
`plt.subplot(412)`  
`plt.plot(trend, label = 'Trend')`  
`plt.legend()`  
  
`plt.subplot(413)`  
`plt.plot(seasonal, label = 'seasonality')`  
`plt.legend()`  
  
`plt.subplot(414)`  
`plt.plot(residual, label = 'Residual')`  
`plt.legend()`  
`plt.show()`



## Making data stationarity

```
In [20]: #Shifting by 1 and Differencing
df['firstDiff'] = df['Sale'] - df['Sale'].shift(1)
df
```

Out[20]:

	Sale	firstDiff
Month		
1964-01-01	2815.0	NaN
1964-02-01	2672.0	-143.0
1964-03-01	2755.0	83.0
1964-04-01	2721.0	-34.0
1964-05-01	2946.0	225.0
...	...	...
1972-05-01	4618.0	-170.0
1972-06-01	5312.0	694.0
1972-07-01	4298.0	-1014.0
1972-08-01	1413.0	-2885.0
1972-09-01	5877.0	4464.0

105 rows × 2 columns

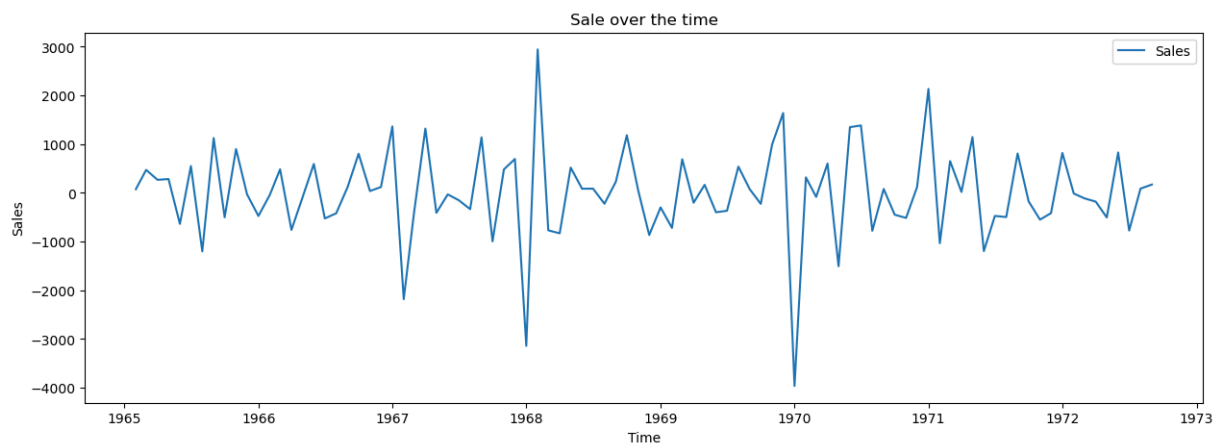
```
In [21]: #Shifting by 12 and Differencing  
df['SeasonalSale'] = df['firstDiff'] - df['firstDiff'].shift(12)  
df
```

Out[21]:

	Sale	firstDiff	SeasonalSale
Month			
1964-01-01	2815.0	NaN	NaN
1964-02-01	2672.0	-143.0	NaN
1964-03-01	2755.0	83.0	NaN
1964-04-01	2721.0	-34.0	NaN
1964-05-01	2946.0	225.0	NaN
...	...	...	...
1972-05-01	4618.0	-170.0	-504.0
1972-06-01	5312.0	694.0	830.0
1972-07-01	4298.0	-1014.0	-773.0
1972-08-01	1413.0	-2885.0	89.0
1972-09-01	5877.0	4464.0	172.0

105 rows × 3 columns

```
In [22]: plt.figure(figsize=(15,5))
plt.title('Sale over the time')
plt.plot(df['SeasonalSale'], label = 'Sales')
plt.xlabel('Time')
plt.ylabel('Sales')
plt.legend()
plt.show()
```



```
In [23]: test_stationarity(df['SeasonalSale'].dropna())
```



#### ADF Test Results

Statistic Value: -4.427713497307538

P-Value: 0.00026504628492931003

Critical Values:

1%: -3.5148692050781247

5%: -2.8984085156250003

10%: -2.58643890625

✔ Series is stationary. Reject the null hypothesis ( $H_0$ ).

In [24]: *#imputing library for ACF and PACF plot*

```
from statsmodels.tsa.stattools import acf, pacf
lag_acf = acf(df['SeasonalSale'].dropna(), nlags=20)
lag_pacf = pacf(df['SeasonalSale'].dropna(), nlags=20)
```

In [25]: lag\_acf

```
Out[25]: array([ 1.          , -0.42195114, -0.09430821,  0.07434433,  0.0043774 ,
        -0.08104485, -0.03665555,  0.04535312,  0.07818965, -0.09834564,
         0.0646181 ,  0.11215009, -0.37966041,  0.3588898 , -0.0527602 ,
        -0.07515121, -0.08502832,  0.06181156,  0.10858334, -0.09818183,
        -0.10793321])
```

In [26]: lag\_pacf

```
Out[26]: array([ 1.          , -0.42658796, -0.34030991, -0.17066773, -0.08834766,
        -0.15160867, -0.22471162, -0.19396218, -0.03684659, -0.11393605,
        -0.03676525,  0.16475062, -0.41289325,  0.09005412,  0.07105284,
         0.04751972, -0.19093674, -0.24796865,  0.01782106, -0.00986552,
        -0.24534446])
```

In [27]: **import** matplotlib.pyplot **as** plt  
**import** statsmodels.api **as** sm

```
plt.figure(figsize=(15, 8))
```

*# ACF plot*

```
plt.subplot(211)
```

```
sm.graphics.tsa.plot_acf(df['SeasonalSale'].dropna(), lags=40, ax=plt.gca())
```


*# PACF plot*

```
plt.subplot(212)
```

```
sm.graphics.tsa.plot_pacf(df['SeasonalSale'].dropna(), lags=40, ax=plt.gca())
```

```
plt.tight_layout()
```

```
plt.show()
```



## ARIMA Model

```
In [28]: from statsmodels.tsa.arima.model import ARIMA
```

```
In [29]: arima_model = ARIMA(df['Sale'], order = (1,1,1)) #p = 1, d = 1, q = 1
         model_fit = arima_model.fit()
```

C:\Users\Admin\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

```
self._init_dates(dates, freq)
```

C:\Users\Admin\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

```
self._init_dates(dates, freq)
```

C:\Users\Admin\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

```
self._init_dates(dates, freq)
```

```
In [30]: model_fit.summary()
```

Out[30]:


SARIMAX Results						
Dep. Variable:		Sale		No. Observations:		105
Model:		ARIMA(1, 1, 1)		Log Likelihood		-952.814
Date:		Fri, 18 Jul 2025		AIC		1911.627
Time:		11:48:43		BIC		1919.560
Sample:		01-01-1964		HQIC		1914.841
		- 09-01-1972				
Covariance Type:		opg				
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.4545	0.114	3.999	0.000	0.232	0.677
ma.L1	-0.9666	0.056	-17.316	0.000	-1.076	-0.857
sigma2	5.226e+06	6.17e+05	8.473	0.000	4.02e+06	6.43e+06
Ljung-Box (L1) (Q):		0.91	Jarque-Bera (JB):		2.59	
Prob(Q):		0.34	Prob(JB):		0.27	
Heteroskedasticity (H):		3.40	Skew:		0.05	
Prob(H) (two-sided):		0.00	Kurtosis:		3.77	

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [31]: df['ArimaPred'] = model_fit.predict(start = 90, end = 105, dynamic = True)
```

```
In [32]: plt.figure(figsize=(15,5))
plt.title('Sale over the time')
plt.plot(df['Sale'], label = 'Actual', color = 'green')
plt.plot(df['ArimaPred'], label = 'ARIMA Predicted', color = 'red')
plt.xlabel('Time')
plt.ylabel('Sales')
plt.legend()
plt.show()
```



As the data is seasonal, So ARIMA model is not giving best prediction

## SARIMAX Model

```
In [33]: import statsmodels.api as sm
```

```
In [34]: model = sm.tsa.statespace.SARIMAX(df['Sale'], order=(1,2,1), seasonal_order=(1,1,1),
      result = model.fit())
```

C:\Users\Admin\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.


```
self._init_dates(dates, freq)
```

C:\Users\Admin\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

```
self._init_dates(dates, freq)
```


```
In [35]: df['forecast'] = result.predict(start = 90, end = 105, dynamic = True)
```

```
In [36]: plt.figure(figsize=(15,5))
plt.title('Sale over the time')
plt.plot(df['Sale'], label = 'Actual', color = 'green')
plt.plot(df['forecast'], label = 'SARIMA Predicted', color = 'red')
plt.xlabel('Time')
plt.ylabel('Sales')
plt.legend()
plt.show()
```



```
In [37]: FuturPred = result.predict(start = 105, end = 120, dynamic = True)
```

```
In [38]: plt.figure(figsize=(15,5))
plt.title('Sale over the time')
plt.plot(df['Sale'], label = 'Actual', color = 'green')
plt.plot(df['forecast'], label = 'SARIMA Predicted', color = 'black')
plt.plot(FuturPred, label = 'Future Prediction', color = 'red')
plt.xlabel('Time')
plt.ylabel('Sales')
plt.legend()
plt.show()
```



```
In [ ]:
```